

# Analysis of IMDB Data

We will analyze a subset of IMDB's actors, genres, movie actors, and movie ratings data. This dataset comes to us from Kaggle (<https://www.kaggle.com/datasets/ashirwadsangwan/imdb-dataset>) although we have taken steps to pull this data into a public s3 bucket:

- s3://cis9760-lecture9-movieanalysis/name.basics.tsv ---> (actors)
- s3://cis9760-lecture9-movieanalysis/title.basics.tsv ---> (genres)
- s3://cis9760-lecture9-movieanalysis/title.principals.tsv ---> (movie actors)
- s3://cis9760-lecture9-movieanalysis/title.ratings.tsv ---> (movie ratings)

## Content

### **name.basics.tsv.gz – Contains the following information for names:**

nconst (string) - alphanumeric unique identifier of the name/person.

primaryName (string)– name by which the person is most often credited.

birthYear – in YYYY format.

deathYear – in YYYY format if applicable, else .

primaryProfession (array of strings)– the top-3 professions of the person.

knownForTitles (array of tconsts) – titles the person is known for.

### **title.basics.tsv.gz - Contains the following information for titles:**

tconst (string) - alphanumeric unique identifier of the title.

titleType (string) – the type/format of the title (e.g. movie, short, tvseries, tvepisode, video, etc).

primaryTitle (string) – the more popular title / the title used by the filmmakers on promotional materials at the point of release.

originalTitle (string) - original title, in the original language.

isAdult (boolean) - 0: non-adult title; 1: adult title.

startYear (YYYY) – represents the release year of a title. In the case of TV Series, it is the series start year.

endYear (YYYY) – TV Series end year. for all other title types.

runtimeMinutes – primary runtime of the title, in minutes.

genres (string array) – includes up to three genres associated with the title.

### **title.principals.tsv – Contains the principal cast/crew for titles:**

tconst (string) - alphanumeric unique identifier of the title.

ordering (integer) – a number to uniquely identify rows for a given title.

nconst (string) - alphanumeric unique identifier of the name/person.

category (string) - the category of job that person was in.

job (string) - the specific job title if applicable, else.

characters (string) - the name of the character played if applicable, else.

### **title.ratings.tsv.gz – Contains the IMDb rating and votes information for titles:**

tconst (string) - alphanumeric unique identifier of the title.

averageRating – weighted average of all the individual user ratings.

numVotes - number of votes the title has received.

## PART 1 - Installation and Initial Setup

Begin by installing the necessary libraries that you may need to conduct your analysis. At the very least, you must install pandas and matplotlib

In [1]: `sc.list_packages()`

VBox()

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
1	application_1668909776641_0002	pyspark	idle	<a href="#">Link</a>	<a href="#">Link</a>	✓

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...

SparkSession available as 'spark'.

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...

Package	Version
-----	-----
beautifulsoup4	4.9.1
boto	2.49.0
click	7.1.2
jmespath	0.10.0
joblib	0.16.0
lxml	4.5.2
mysqlclient	1.4.2
nlTK	3.5
nose	1.3.4
numpy	1.16.5
pip	9.0.1
py-dateutil	2.2
python37-sagemaker-pyspark	1.4.0
pytz	2020.1
PyYAML	5.3.1
regex	2020.7.14
setuptools	28.8.0
six	1.13.0
soupsieve	1.9.5
tqdm	4.48.2
wheel	0.29.0
windmill	1.6

Let's install the necessary packages here

```
In [2]: sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...  
Collecting pandas==1.0.3  
Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1\_x86\_64.whl  
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)  
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)  
Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)  
Using cached https://files.pythonhosted.org/packages/36/7a/87837f39d0296e723bb9b62bb257d0355c7f6128853c78955f57342a56d/python\_dateutil-2.8.2-py2.py3-none-any.whl  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)  
Installing collected packages: python-dateutil, pandas  
Successfully installed pandas-1.0.3 python-dateutil-2.8.2

Collecting matplotlib==3.2.1  
Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3ba1f09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1\_x86\_64.whl  
Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1668917922725-0/lib/python3.7/site-packages (from matplotlib==3.2.1)  
Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)  
Using cached https://files.pythonhosted.org/packages/6c/10/a7d0fa5baea8fe7b50f448ab742f26f52b80bfca85ac2be9d35cdd9a3246/pyparsing-3.0.9-py3-none-any.whl  
Collecting cyclor>=0.10 (from matplotlib==3.2.1)  
Using cached https://files.pythonhosted.org/packages/5c/f9/695d6bedebd747e5eb0fe8fad57b72fdf25411273a39791cde838d5a8f51/cyclor-0.11.0-py3-none-any.whl  
Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)  
Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)  
Using cached https://files.pythonhosted.org/packages/ab/8f/8dbe2d4efc4c0b08ec67d6efb7cc31fbfd688c80afad85f65980633b0d37/kiwisolver-1.4.4-cp37-cp37m-manylinux\_2\_5\_x86\_64.manylinux1\_x86\_64.whl  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)  
Collecting typing-extensions; python\_version < "3.8" (from kiwisolver>=1.0.1->matplotlib==3.2.1)  
Using cached https://files.pythonhosted.org/packages/0b/8e/f1a0a5a76cfef77e1eb6004cb49e5f8d72634da638420b9ea492ce8305e8/typing\_extensions-4.4.0-py3-none-any.whl  
Installing collected packages: pyparsing, cyclor, typing-extensions, kiwisolver, matplotlib  
Successfully installed cyclor-0.11.0 kiwisolver-1.4.4 matplotlib-3.2.1 pyparsing-3.0.9 typing-extensions-4.4.0

Now, import the installed packages from the previous block below.

```
In [3]: import numpy as np
import matplotlib.pyplot as plt
```

```
from pyspark.sql.functions import approx_count_distinct, avg, collect_set, countDistinct
```

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

## Loading Data

Load all data from S3 into a Spark dataframe object

```
In [4]: actors = spark.read.csv('s3://cis9760-lecture9-movieanalysis/name.basics.tsv', sep=r'\n')
genres = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.basics.tsv', sep=r'\n')
movie_actors = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.principals.tsv', sep=r'\n')
movie_ratings = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.ratings.tsv', sep=r'\n')
```

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

## Actors

Display the schema below:

```
In [5]: actors.printSchema()
```

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
root
```

```
|-- nconst: string (nullable = true)
|-- primaryName: string (nullable = true)
|-- birthYear: string (nullable = true)
|-- deathYear: string (nullable = true)
|-- primaryProfession: string (nullable = true)
|-- knownForTitles: string (nullable = true)
```

Display the first 5 rows with the following columns:

- primaryName
- birthYear
- deathYear
- knownForTitles

```
In [6]: actors.select("primaryName", "birthYear", "deathYear", "knownForTitles").show(5)
```

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
+-----+-----+-----+-----+
| primaryName|birthYear|deathYear|      knownForTitles|
+-----+-----+-----+-----+
| Fred Astaire|    1899|    1987|tt0050419,tt00531...|
| Lauren Bacall|    1924|    2014|tt0071877,tt01170...|
|Brigitte Bardot|    1934|      \N|tt0054452,tt00491...|
| John Belushi|    1949|    1982|tt0077975,tt00725...|
| Ingmar Bergman|    1918|    2007|tt0069467,tt00509...|
+-----+-----+-----+-----+
only showing top 5 rows
```

## Genres

Display the first 10 rows with the following columns:

- titleType
- primaryTitle
- genres

```
In [7]: genres.select("titleType", "primaryTitle", "genres").show(10)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+-----+-----+-----+
|titleType|      primaryTitle|      genres|
+-----+-----+-----+-----+
| short| Carmencita| Documentary,Short|
| short|Le clown et ses c...| Animation,Short|
| short| Pauvre Pierrot|Animation,Comedy,...|
| short| Un bon bock| Animation,Short|
| short| Blacksmith Scene| Comedy,Short|
| short| Chinese Opium Den| Short|
| short|Corbett and Court...| Short,Sport|
| short|Edison Kinetoscop...| Documentary,Short|
| movie| Miss Jerry| Romance|
| short| Exiting the Factory| Documentary,Short|
+-----+-----+-----+-----+
only showing top 10 rows
```

Display the unique categories below:

```
In [8]: genres.select('titleType').distinct().show()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
```

```
+-----+
|  titleType|
+-----+
|  tvSeries|
|tvMiniSeries|
|      movie|
|  videoGame|
|  tvSpecial|
|      video|
|      tvMovie|
|  tvEpisode|
|      tvShort|
|      short|
+-----+
```

Display the schema below:

In [9]: `genres.printSchema()`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
root
|-- tconst: string (nullable = true)
|-- titleType: string (nullable = true)
|-- primaryTitle: string (nullable = true)
|-- originalTitle: string (nullable = true)
|-- isAdult: string (nullable = true)
|-- startYear: string (nullable = true)
|-- endYear: string (nullable = true)
|-- runtimeMinutes: string (nullable = true)
|-- genres: string (nullable = true)
```

## Movie Actors

Display the schema below:

In [10]: `movie_actors.printSchema()`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
root
|-- tconst: string (nullable = true)
|-- ordering: string (nullable = true)
|-- nconst: string (nullable = true)
|-- category: string (nullable = true)
|-- job: string (nullable = true)
|-- characters: string (nullable = true)
```

Display the first 10 rows below

In [11]: `movie_actors.show(10)`

```
VBox()
```

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
```

tconst	ordering	nconst	category	job	characters
tt0000001	1	nm1588970	self	\N	["Herself"]
tt0000001	2	nm0005690	director	\N	\N
tt0000001	3	nm0374658	cinematographer	director of photo...	\N
tt0000002	1	nm0721526	director	\N	\N
tt0000002	2	nm1335271	composer	\N	\N
tt0000003	1	nm0721526	director	\N	\N
tt0000003	2	nm5442194	producer	producer	\N
tt0000003	3	nm1335271	composer	\N	\N
tt0000003	4	nm5442200	editor	\N	\N
tt0000004	1	nm0721526	director	\N	\N

only showing top 10 rows

## Movie Ratings

Display the schema below:

```
In [12]: movie_ratings.printSchema()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
root
|-- tconst: string (nullable = true)
|-- averageRating: string (nullable = true)
|-- numVotes: string (nullable = true)
```

Display the first 10 rows in a descending order by the number of votes

```
In [13]: movie_ratings.orderBy(col('numVotes').desc()).show(10)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+-----+-----+
| tconst|averageRating|numVotes|
+-----+-----+-----+
|tt7430722|        6.8|    9999|
|tt4445154|        8.1|    9997|
|tt2229907|        6.3|    9996|
|tt0294097|        8.0|    9994|
|tt0264734|        6.5|    9993|
|tt8860450|        6.3|    9991|
|tt2032572|        5.2|    9991|
|tt0664505|        8.4|     999|
|tt7508752|        7.9|     999|
|tt1077089|        7.3|     999|
+-----+-----+-----+
only showing top 10 rows
```

# Overview of Data

Display the number of rows and columns in each DataFrame object.

```
In [14]: actors_cols = len(actors.columns)
print(f"Number of columns in Actors table: {actors_cols}")
actors_rows = actors.count()
print(f"Number of rows in Actors table: {actors_rows}", "\n")

genres_cols = len(genres.columns)
print(f"Number of columns in Genres table: {genres_cols}")
genres_rows = genres.count()
print(f"Number of rows in Genres table: {genres_rows}", "\n")

movieactors_cols = len(movie_actors.columns)
print(f"Number of columns in Movie Actors table: {movieactors_cols}")
movieactors_rows = movie_actors.count()
print(f"Number of rows in Movie Actors table: {movieactors_rows}", "\n")

movieratings_cols = len(movie_ratings.columns)
print(f"Number of columns in Movie Ratings table: {movieratings_cols}")
movieratings_rows = movie_ratings.count()
print(f"Number of rows in Movie Ratings table: {movieratings_rows}", "\n")
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...  
Number of columns in Actors table: 6  
Number of rows in Actors table: 9706922

Number of columns in Genres table: 9  
Number of rows in Genres table: 6321302

Number of columns in Movie Actors table: 6  
Number of rows in Movie Actors table: 36468817

Number of columns in Movie Ratings table: 3  
Number of rows in Movie Ratings table: 993153

## PART 2 - Analyzing Genres

Let's now answer this question: how many unique genres are represented in this dataset?

Essentially, we have the genres per movie as a list - this is useful to quickly see what each movie might be represented as but it is difficult to easily answer questions such as:

- How many movies are categorized as Comedy, for instance?
- What are the top 20 most popular genres available?

## Association Table



We need to "break out" these genres from the tconst? One common approach to take is to build an association table mapping a single tconst multiple times to each distinct genre.

For instance, given the following:

tconst	titleType	genres
abcd123	XXX	a,b,c

We would like to derive something like:

tconst	titleType	genre
abcd123	XXX	a
abcd123	XXX	b
abcd123	XXX	c

What this does is allow us to then perform a myriad of rollups and other analysis on this association table which can aid us in answering the questions asked above.

Implement the code necessary to derive the table described from the data set

```
In [15]: genres.select('tconst','titleType','genres').show(5)
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...

```
+-----+
|  tconst|titleType|          genres|
+-----+
|tt0000001|  short| Documentary,Short|
|tt0000002|  short|  Animation,Short|
|tt0000003|  short|Animation,Comedy,...|
|tt0000004|  short|  Animation,Short|
|tt0000005|  short|    Comedy,Short|
+-----+
```

only showing top 5 rows

Display the first 10 rows of your association table below

```
In [16]: association_table = genres.select("tconst","titleType",split(col("genres"),",").alias("genres_split"))
association_table2=association_table.select("tconst","titleType",explode(association_table.genres_split).alias("genre")).show(10)
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...

tconst	titleType	genres
tt0000001	short	Documentary
tt0000001	short	Short
tt0000002	short	Animation
tt0000002	short	Short
tt0000003	short	Animation
tt0000003	short	Comedy
tt0000003	short	Romance
tt0000004	short	Animation
tt0000004	short	Short
tt0000005	short	Comedy

only showing top 10 rows

## Total Unique Genres

What is the total number of unique genres available in the movie category?

```
In [17]: association_count = association_table2.select(countDistinct("genres"))
         association_count.show()

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+
|count(DISTINCT genres)|
+-----+
|                        29|
+-----+
```

What are the unique genres available?

```
In [18]: association_table2.select('genres').distinct().show()

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
```

```

+-----+
|   genres|
+-----+
|   Mystery|
|   Musical|
|   Sport|
|   Action|
| Talk-Show|
|   Romance|
|   Thriller|
|      \N|
| Reality-TV|
|   Family|
|   Fantasy|
|   History|
| Animation|
| Film-Noir|
|   Short|
|   Sci-Fi|
|   News|
|   Drama|
| Documentary|
|   Western|
+-----+

```

only showing top 20 rows

**Oops! Something is off!**

```

In [19]: nll = '\\N'
association_table3 = association_table2.select('genres').filter(col('genres') != nll).
association_table3.show()

```

VBox()

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

genres
Mystery
Musical
Sport
Action
Talk-Show
Romance
Thriller
Reality-TV
Family
Fantasy
History
Animation
Short
Film-Noir
Sci-Fi
News
Drama
Documentary
Western
Comedy

only showing top 20 rows

```
In [20]: association_table3.select(countDistinct("genres")).show()

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+
|count(DISTINCT genres)|
+-----+
|                        28|
+-----+
```

# Top Genres by Movies

Now let's find the highest rated genres in this dataset by rolling up genres.

## Average Rating / Genre

So now, let's unroll our distinct count a bit and display the per average rating value of per genre.

The expected output should be:

genre	averageRating
a	8.5
b	6.3
c	7.2

Or something to that effect.

First, let's join our two dataframes (movie ratings and genres) by tconst

```
In [21]: nll = '\\N'
association_table4 = association_table2.filter(col('genres') != nll)
top_genres = association_table4.join(movie_ratings).where(association_table4['tconst'])
top_genres2 = top_genres.select('genres', 'averageRating')
top_genres2.show(10)
```

VBox()

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
+-----+-----+
|   genres|averageRating|
+-----+-----+
|Documentary|          5.1|
|      Short|          5.1|
|Documentary|          5.2|
|      Short|          5.2|
|      Comedy|         5.2|
|      Short|          5.2|
|      Comedy|         6.0|
|      Horror|         6.0|
|      Short|         6.0|
|Documentary|          4.9|
+-----+-----+
```

only showing top 10 rows

Now, let's aggregate along the averageRating column to get a resultant dataframe that displays average rating per genre.

```
In [22]: top_genres3 = top_genres2.withColumn('averageRating', col('averageRating').cast('Float'))
top_genres4 = top_genres3.groupBy('genres').agg(avg("averageRating").alias("averageRating"))
top_genres4.show()
```

VBox()

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
+-----+-----+
|   genres|   averageRating|
+-----+-----+
|   Mystery| 7.215679898056961|
|   Musical| 6.544660195495522|
|   Action| 6.951029447217718|
|   Sport| 6.995047307520108|
| Talk-Show| 6.598412160498378|
|   Romance| 6.784248168750174|
|   Thriller| 6.312686083353854|
| Reality-TV| 6.8388670073012765|
|   Family| 6.989731263527516|
|   Fantasy| 7.093731237046438|
|   History| 7.304633665280264|
| Animation| 7.046786062925208|
| Film-Noir| 6.636246780503378|
|   Short| 6.799363502678704|
|   Sci-Fi| 6.747496253901592|
|   News| 6.467539497212473|
|   Drama| 7.018453637663789|
| Documentary| 7.241740584659479|
|   Western| 7.109783420299982|
|   Comedy| 6.919198969550161|
+-----+-----+
only showing top 20 rows
```

## Horizontal Bar Chart of Top Genres

With this data available, let us now build a barchart of all genres

**HINT:** don't forget about the matplotlib magic!

```
%matplotlib plt
```

```
In [23]: top_genres4.sort(col('averageRating').desc()).show()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
```

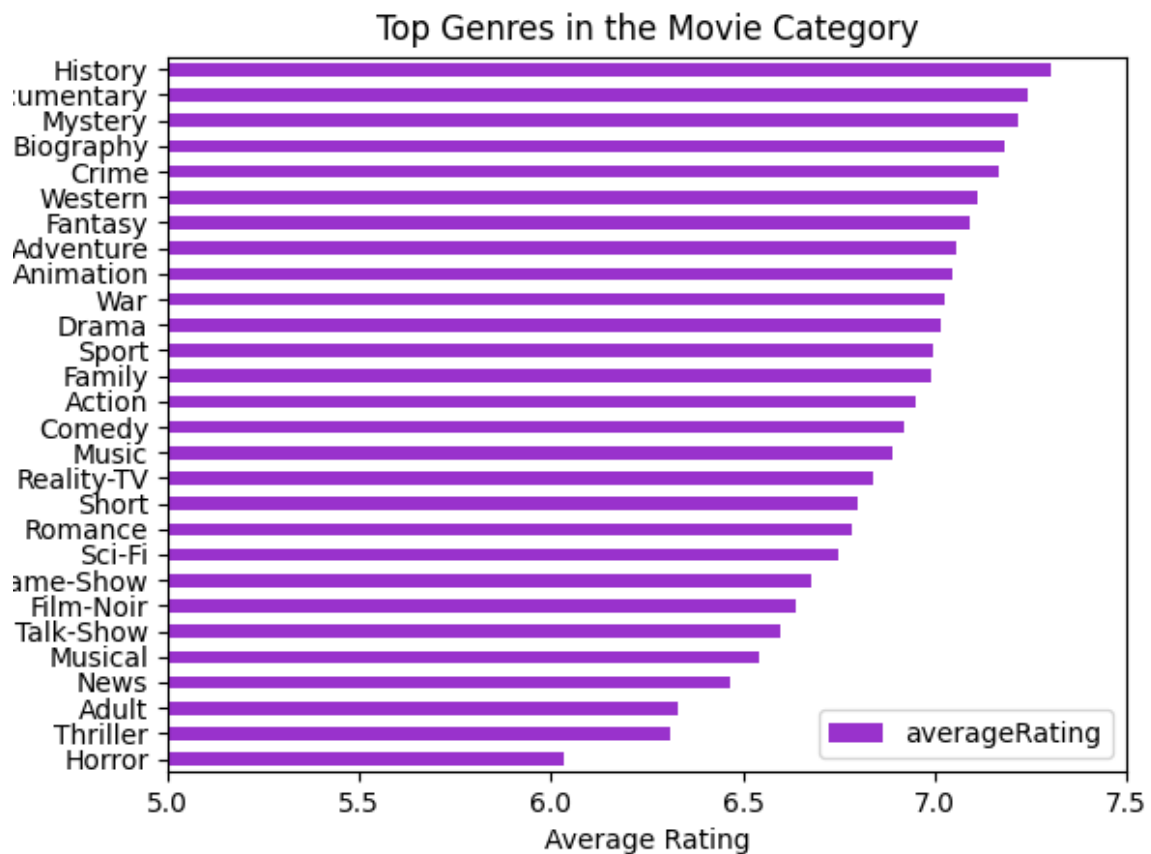
genres	averageRating
History	7.304633665280264
Documentary	7.241740584659479
Mystery	7.215679898056961
Biography	7.180114836990299
Crime	7.165008360689229
Western	7.109783420299982
Fantasy	7.093731237046438
Adventure	7.0567301815847685
Animation	7.046786062925208
War	7.026155360390018
Drama	7.018453637663789
Sport	6.995047307520108
Family	6.989731263527516
Action	6.951029447217718
Comedy	6.919198969550161
Music	6.890572244550894
Reality-TV	6.8388670073012765
Short	6.799363502678704
Romance	6.784248168750174
Sci-Fi	6.747496253901592

only showing top 20 rows

```
In [24]: hbchart=top_genres4.select('genres','averageRating').sort(col('averageRating')).toPandas
hbchart.plot.barh(x='genres', y='averageRating', color='darkorchid')
plt.xlim([5.0,7.5])
plt.xlabel('Average Rating')
plt.ylabel('Genre')
plt.title('Top Genres in the Movie Category')
plt.legend(loc=4)
%matplotlib plt
```

VBox()

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...



## PART 3 - Analyzing Job Categories

### Total Unique Job Categories

**What is the total number of unique job categories?**

```
In [25]: movie_actors.select('tconst', 'category').show(5)
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...

tconst	category
tt0000001	self
tt0000001	director
tt0000001	cinematographer
tt0000002	director
tt0000002	composer

only showing top 5 rows

```
In [26]: movie_actors.select(countDistinct('category')).show()
```

VBox()



```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+
|count(DISTINCT category)|
+-----+
|                          12|
+-----+
```

### What are the unique job categories available?

In [27]: `movie_actors.select('category').distinct().show()`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+
|          category|
+-----+
|          actress|
|          producer|
|production_designer|
|          writer|
|          actor|
| cinematographer|
|  archive_sound|
|  archive_footage|
|          self|
|          editor|
|          composer|
|          director|
+-----+
```

## Top Job Categories

Now let's find the top job categories in this dataset by rolling up categories.

### Counts of Titles / Job Category

The expected output should be:

category	count
a	15
b	2
c	45

Or something to that effect.

In [28]: `movie_actors.groupBy('category').count().show()`

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+
|          category|
+-----+
|          actress|
|          producer|
|production_designer|
|          writer|
|          actor|
| cinematographer|
|  archive_sound|
|  archive_footage|
|          self|
|          editor|
|          composer|
|          director|
+-----+
```

```
+-----+-----+
|          category|  count|
+-----+-----+
|          actress|6325097|
|          producer|2197866|
|production_designer| 285924|
|          writer|4811596|
|          actor|8493701|
|    cinematographer|1300404|
|    archive_sound|   2143|
|    archive_footage| 209035|
|          self|6153089|
|          editor|1197669|
|        composer|1313187|
|        director|4179106|
+-----+-----+
```

## Bar Chart of Top Job Categories

With this data available, let us now build a barchart of the top 5 categories.

**HINT:** don't forget about the matplotlib magic!


```
%matplotlib plt
```

```
In [29]: movie_actors2 = movie_actors.groupby('category').count().sort(col('count').desc())
movie_actors2.show()

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+-----+
|          category|  count|
+-----+-----+
|          actor|8493701|
|          actress|6325097|
|          self|6153089|
|          writer|4811596|
|          director|4179106|
|          producer|2197866|
|          composer|1313187|
|    cinematographer|1300404|
|          editor|1197669|
|production_designer| 285924|
|    archive_footage| 209035|
|    archive_sound|   2143|
+-----+-----+
```

```
In [30]: bctjc=movie_actors2.limit(5).toPandas()
bctjc.plot.bar(x='category', y='count', color='orange')
plt.ylim([3000000, 9000000])
plt.xlabel('Job Categories')
plt.ylabel('Count')
plt.title('Top Job Categories')
plt.legend(loc=1)
%matplotlib plt
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
```



## PART 4 - Answer to the following questions:

### 1) Find all the "movies" featuring "Johnny Depp" and "Helena Bonham Carter".

First join actors, genres, and movie actors on each other

```
In [31]: jd_hbc = actors.join(movie_actors).where(actors['nconst'] == movie_actors['nconst'])
jd_hbc2 = jd_hbc.join(genres).where(jd_hbc['tconst'] == genres['tconst']).select('prim
jd_hbc3 = jd_hbc2.select('primaryTitle', 'primaryName', 'titleType').filter("titleType =
jd_hbc4 = jd_hbc3.select('primaryTitle', 'primaryName').filter("primaryName=='Johnny De
jd_hbc4.groupBy('primaryTitle').count().where("count>1").drop('count').show(truncate=F
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
```

```
+-----+
|primaryTitle|
+-----+
|Corpse Bride|
|Dark Shadows|
|Charlie and the Chocolate Factory|
|Alice Through the Looking Glass|
|Sweeney Todd: The Demon Barber of Fleet Street|
|Alice in Wonderland|
+-----+
```

## 2) Find all the "movies" featuring "Brad Pitt" after 2010.

```
In [32]: nll = '\\N'
bp = actors.join(movie_actors).where(actors['nconst'] == movie_actors['nconst'])
bp2 = bp.join(genres).where(bp['tconst'] == genres['tconst']).select('primaryTitle', '
bp2.select('primaryTitle', 'startYear').sort(col('startYear').desc()).show(truncate=False)
```

VBox()

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
+-----+-----+
|primaryTitle|startYear|
+-----+-----+
|Babylon|2021|
|Kajillionaire|2020|
|Irresistible|2020|
|Ad Astra|2019|
|Once Upon a Time ... in Hollywood|2019|
|The King|2019|
|Vice|2018|
|War Machine|2017|
|Allied|2016|
|Voyage of Time: Life's Journey|2016|
|The Big Short|2015|
|By the Sea|2015|
|Hitting the Apex|2015|
|Fury|2014|
|World War Z|2013|
|Kick-Ass 2|2013|
|12 Years a Slave|2013|
|Killing Them Softly|2012|
|The Tree of Life|2011|
|Moneyball|2011|
+-----+-----+
```

## 3) What is the number of "movies" "acted" by "Zendaya" per year?

```
In [33]: nll = '\\N'
zendaya = actors.join(movie_actors).where(actors['nconst'] == movie_actors['nconst'])
zendaya2 = zendaya.join(genres).where(zendaya['tconst'] == genres['tconst']).select('p
zendaya2.groupBy('startYear').count().show(truncate=False)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+-----+
|startYear|count|
+-----+-----+
|2020      |1    |
|2018      |2    |
|2017      |1    |
+-----+-----+
```

#### 4) What are the "movies" by average rating greater than "9.7" and released in "2019"?

```
In [34]: movies = genres.join(movie_ratings).where(genres['tconst'] == movie_ratings['tconst'])
movies2 = movies.select('primaryTitle', 'averageRating').filter((movies.titleType=='mc
movies3 = movies2.withColumn('averageRating', col('averageRating').cast('Float'))
movies3.sort(col('averageRating').desc()).show(truncate=False)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(hei
ght='25px', width='50%'),...
+-----+-----+
|primaryTitle                                     |averageRating|
+-----+-----+
|Our Scripted Life                               |10.0         |
|The Twilight Zone: A 60th Anniversary Celebration|10.0         |
|A Grunt's Life                                  |10.0         |
|Bu Can Var Oldugu Sürece                       |10.0         |
|L'Enfant Terrible                              |10.0         |
|Kirket                                           |10.0         |
|A Medicine for the Mind                        |10.0         |
|The Butcher Baronet                            |10.0         |
|Love in Kilnerry                               |10.0         |
|Puritan: All of Life to The Glory of God        |9.9          |
|The Cardinal                                    |9.9          |
|Superhombre                                    |9.9          |
|Kamen Rider Zi-O: Over Quartzer                |9.8          |
|Square One                                     |9.8          |
|Time and motion                               |9.8          |
|We Shall Not Die Now                           |9.8          |
|From Shock to Awe                              |9.8          |
|Randhawa                                       |9.8          |
|Gini Helida Kathe                             |9.8          |
+-----+-----+
```

### Extra Credit - Analysis of your choice

Try and analyze some interesting dimension to this data. You should specify the question in your Project2\_Analysis.ipynb.

You must join at least two datasets.

# Top 20 Movie Directors with the highest Weighted Average Ratings released in 2017-2022.

```
In [35]: nll = '\\N'
actor = genres.join(movie_actors).where(genres['tconst'] == movie_actors['tconst'])
actor2 = actor.join(actors).where(actor['nconst'] == actors['nconst']).select('primary
actor3 = actor2.join(movie_ratings).where(actor2['tconst'] == movie_ratings['tconst'])
actor4 = actor3.filter((genres.titleType=='movie') & (genres.startYear != nll) & (genres.s
actor5 = actor4.withColumn('averageRating', col('averageRating').cast('Float')).withCo
actor6 = actor5.groupby('primaryName').sum('Weighted Score').withColumnRenamed('sum(We
```

VBox()

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(hei  
ght='25px', width='50%'),...

primaryName	Weighted Score
Joe Russo	1.1287869E7
Anthony Russo	1.1286742E7
Jon Watts	5174425.25
David Leitch	5025152.625
James Mangold	4872498.5
Jordan Peele	4410769.125
Todd Phillips	4108825.75
Taika Waititi	4096711.0
Ryan Coogler	4024088.5
James Gunn	3883387.25
Chad Stahelski	3872830.125
Christopher Nolan	3851313.25
Patty Jenkins	3753087.75
Andy Muschietti	3746174.5
Rian Johnson	3439359.900390625
Steven Spielberg	3220537.5
Denis Villeneuve	3199704.0
Bryan Singer	3140792.0
Martin McDonagh	3016148.5
Edgar Wright	2941025.25

only showing top 20 rows